

# Masinsko ucenje u R-u

Igor Hut

December 5, 2016

## Contents

Uvodna analiza podataka . . . . .	1
Primer inicijalne provere i analize podataka . . . . .	1
Regresija, klasifikacija, klasterizacija - Uvod . . . . .	4
Linearna regresija - uvodni primer 1 . . . . .	4
Regresija - uvodni primer 2 . . . . .	5
Klasifikacija - uvodni primer . . . . .	5
Klasterovanje - uvodni primer . . . . .	6
Ocena modela . . . . .	8
Konfuzionna matrica - Primeri . . . . .	8
Primer 1: . . . . .	8
Primer 2: . . . . .	10
Tacnost, preciznost, senzitivnost (recall), specificnost - Primer . . . . .	11
Zadatak za vezbanje na casu: . . . . .	12
Kvalitet regresije . . . . .	12
Primer: . . . . .	12
Procena valjanosti klasterizacije: WSS vs BSS . . . . .	13
Trening set i test set . . . . .	15
Primer . . . . .	15
Zadatak za vezbanje na casu: . . . . .	17
Upotreba unakrsne validacije (cross-validation) . . . . .	17
Bajas i varijansa (Bias and Variance) . . . . .	17
Primer . . . . .	17
Regresija . . . . .	20
Prosta linearna regresija . . . . .	20
Koriscenje funkcije <code>lm()</code> - Primeri . . . . .	20
Multivarijabilna linearna regresija . . . . .	25
Primer 1 . . . . .	25
Primer 2 . . . . .	31

## Uvodna analiza podataka

- `data.frame` - primarna forma za smestanje, analizu i manipulaciju podacima
- Uvoz podataka - licna preporuka `readr` (Hadley Wickham)
- Prvi korak - upoznavanje sa podacima
- Dimenzije - broj promenljivih (features) i opservacija (observations)
- Ciscenje i uredjivanje podataka - licna preporuka `dplyr` i `tidyr` (Hadley Wickham)
- Vizualizacija - `base` i `ggplot2` (Hadley Wickham)

## Primer inicijalne provere i analize podataka

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 3.3.2
```

```

library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(tidyr)
# broj promenljivih i broj opservacija
str(iris)

## 'data.frame':   150 obs. of  5 variables:
## $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

dim(iris)

## [1] 150  5

# Nekoliko prvih i poslednjih vrsta iz `iris` baze
head(iris)

##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1          5.1          3.5          1.4          0.2  setosa
## 2          4.9          3.0          1.4          0.2  setosa
## 3          4.7          3.2          1.3          0.2  setosa
## 4          4.6          3.1          1.5          0.2  setosa
## 5          5.0          3.6          1.4          0.2  setosa
## 6          5.4          3.9          1.7          0.4  setosa

tail(iris)

##   Sepal.Length Sepal.Width Petal.Length Petal.Width  Species
## 145          6.7          3.3          5.7          2.5 virginica
## 146          6.7          3.0          5.2          2.3 virginica
## 147          6.3          2.5          5.0          1.9 virginica
## 148          6.5          3.0          5.2          2.0 virginica
## 149          6.2          3.4          5.4          2.3 virginica
## 150          5.9          3.0          5.1          1.8 virginica

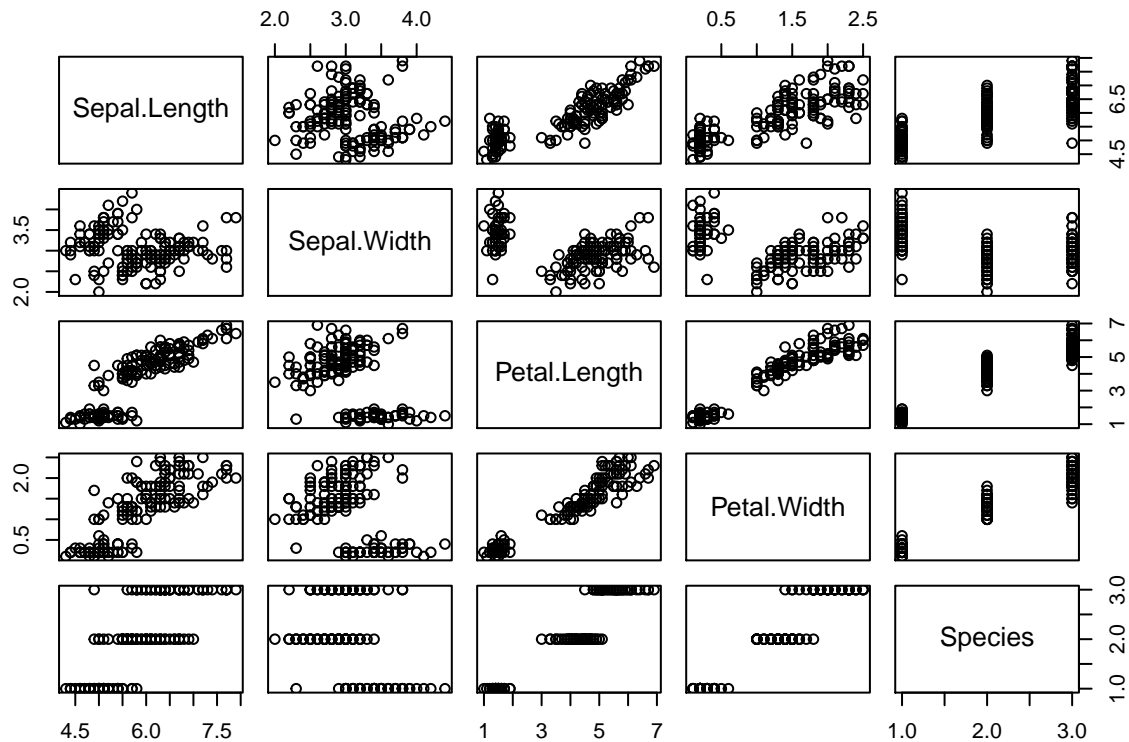
# sumarna statistika za podatke u `iris` bazi
summary(iris)

##   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width
## Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100
## 1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300
## Median :5.800   Median :3.000   Median :4.350   Median :1.300
## Mean   :5.843   Mean   :3.057   Mean   :3.758   Mean   :1.199
## 3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
## Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500


```

```
##      Species
## setosa   :50
## versicolor:50
## virginica :50
##
##
##
```

`plot(iris)` *#rezultat ce biti isti kao da smo upotrebili 'pairs' funkciju iz 'base' biblioteke*



```
ggplot(iris, aes( x = Sepal.Width, y = Sepal.Length, col = Species)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE)
```



## Regresija, klasifikacija, klasterizacija - Uvod

### Linearna regresija - uvodni primer 1

*# Ucitavamo "Wage" bazu iz "ISLR" paketa koja sadrzi neke opste podatke o radnicima u srednje-Atlanstkom okeanu*  
`library(ISLR)`

## Warning: package 'ISLR' was built under R version 3.3.2

`data("Wage")`

*# Generisemo linearni model "lm\_wage"*

`lm_wage <- lm(wage ~ age, data = Wage)`

*# ?lm - kako se funkcija "lm()" koristi*

*# Definisemo data.frame sa novim vrednostima, koje nisu koriscenje za sintezu modela: "unseen"*  
`unseen <- data.frame(age = 60)`

*# Na osnovu modela "lm\_wage" predvidjamo koliko iznosi plata 60-togodisnjeg radnika*  
`predict(lm_wage, unseen)`

## 1

```
## 124.1413
```

## Regresija - uvodni primer 2

```
# Broj pregleda vased LinkedIn profila u periodu od tri nedelje
linkedin <- c(5, 7, 4, 9, 11, 10, 14, 17, 13, 11, 18, 17, 21, 21, 24, 23, 28, 35, 21, 27, 23)


# Vektor koji sadrzi korespodentne dane: "dani"
dani <- 1:21

# Linearni model - broj pregleda po danima: linkedin_lm

linkedin_lm <- lm(linkedin ~ dani)

# Predviđamo broj pregleda u sledeca tri dana: linkedin_pred
buduci_dani <- data.frame(dani = 22:24)
linkedin_pred <- predict(linkedin_lm, buduci_dani)

# Plotujemo "istorijske" podatke i predviđanje
plot(linkedin ~ dani, xlim = c(1, 24))
points(22:24, linkedin_pred, col = "red", pch = 16)
```



## Klasifikacija - uvodni primer

Primer neadekvatnog klasifikatora - krajnje overfitovanje podataka iz trening seta!

```

library(readr)

## Warning: package 'readr' was built under R version 3.3.2
if (!"emails" %in% ls()) {
  emails <- read_csv("data/emails_small.csv")
}

## Parsed with column specification:
## cols(
##   avg_capital_seq = col_double(),
##   spam = col_integer()
## )

# Proveravamo strukturu seta podataka
str(emails)

## Classes 'tbl_df', 'tbl' and 'data.frame':   13 obs. of  2 variables:
## $ avg_capital_seq: num  1 2.11 4.12 1.86 2.97 ...
## $ spam           : int  0 0 1 0 1 0 1 0 0 1 ...
## - attr(*, "spec")=List of 2
## ..$ cols      :List of 2
## .. ..$ avg_capital_seq: list()
## .. ..$- attr(*, "class")= chr  "collector_double" "collector"
## .. ..$ spam       : list()
## .. ..$- attr(*, "class")= chr  "collector_integer" "collector"
## ..$ default: list()
## .. ..$- attr(*, "class")= chr  "collector_guess" "collector"
## ..$- attr(*, "class")= chr "col_spec"

# Definisemo funkciju spam_classifier()
# 1 - spam, 0 - ham
spam_classifier <- function(x){
  prediction <- rep(NA,length(x))
  prediction[x > 4] <- 1
  prediction[x >= 3 & x <= 4] <- 0
  prediction[x >= 2.2 & x < 3] <- 1
  prediction[x >= 1.4 & x < 2.2] <- 0
  prediction[x > 1.25 & x < 1.4] <- 1
  prediction[x <= 1.25] <- 0
  return(prediction)
}

# Primenimo nas klasifikator na kolonu "avg_capital_seq": "spam_pred"
spam_pred <- spam_classifier(emails$avg_capital_seq)

# Uporedimo "spam_pred" i "emails$spam"
spam_pred == emails$spam

## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
identical(spam_pred, as.numeric(emails$spam))

## [1] TRUE

```

Klasterovanje - uvodni primer

```

# Da bi smo obezbedili reproduktivnost
set.seed(1)

# Proveravamo strukturu podataka
str(iris)

## 'data.frame':  150 obs. of  5 variables:
## $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

head(iris)

##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1         5.1         3.5         1.4         0.2  setosa
## 2         4.9         3.0         1.4         0.2  setosa
## 3         4.7         3.2         1.3         0.2  setosa
## 4         4.6         3.1         1.5         0.2  setosa
## 5         5.0         3.6         1.4         0.2  setosa
## 6         5.4         3.9         1.7         0.4  setosa

# Delimo "iris" na dva seta: "my_iris" i "species"
my_iris <- iris[-5]
species <- iris$Species



# Vrsimo k-means klasterizaciju za "my_iris", pretpostavljamo da postoje tri klase: "kmeans_iris"
kmeans_iris <- kmeans(my_iris,3)

# Poredimo dobijene klasterne sa istinskim klasama (kategorijama)
table(species, kmeans_iris$cluster)

##
## species      1  2  3
## setosa      50  0  0
## versicolor  0  2 48
## virginica   0 36 14

# Plotujemo "Petal.Width" vs "Petal.Length", bojimo po klasterima odn. postojećim kategorijama
par(mfrow = c(1,2))
plot(Petal.Length ~ Petal.Width, data = my_iris, col = kmeans_iris$cluster)
title("k-means - klasteri")
plot(Petal.Length ~ Petal.Width, data = my_iris, col = iris$Species)
title("Istinske klase")

```



## Ocena modela

### Konfuziona matrica - Primeri

#### Primer 1:

```
library(rpart)
library(readr)
library(purrr)
```

```
##
## Attaching package: 'purrr'
## The following objects are masked from 'package:dplyr':
##
##   contains, order_by
```

```
# Import podataka
if (!"titanic" %in% ls()) {
  titanic <- read_csv("data/train.csv")
}
```

```
## Parsed with column specification:
## cols(
##   survived = col_integer(),
##   pclass = col_integer(),
##   name = col_character(),
```



```
## sex = col_character(),
## age = col_double(),
## sibsp = col_integer(),
## parch = col_integer(),
## ticket = col_character(),
## fare = col_double(),
## cabin = col_character(),
## embarked = col_character()
## )
```

```
# Da obezbedimo reproduktivnost
set.seed(33)
```

```
# Proveravamo strukturu data seta
str(titanic)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame': 891 obs. of 11 variables:
## $ survived: int 0 1 1 1 0 0 0 0 1 1 ...
## $ pclass : int 3 1 3 1 3 3 1 3 3 2 ...
## $ name : chr "Braund, Mr. Owen Harris" "Cumings, Mrs. John Bradley (Florence Briggs Thayer)" "H
## $ sex : chr "male" "female" "female" "female" ...
## $ age : num 22 38 26 35 35 NA 54 2 27 14 ...
## $ sibsp : int 1 1 0 1 0 0 0 3 0 1 ...
## $ parch : int 0 0 0 0 0 0 0 1 2 0 ...
## $ ticket : chr "A/5 21171" "PC 17599" "STON/O2. 3101282" "113803" ...
## $ fare : num 7.25 71.28 7.92 53.1 8.05 ...
## $ cabin : chr NA "C85" NA "C123" ...
## $ embarked: chr "S" "C" "S" "S" ...
## - attr(*, "spec")=List of 2
## ..$ cols :List of 11
## .. ..$ survived: list()
## .. .. ..- attr(*, "class")= chr "collector_integer" "collector"
## .. ..$ pclass : list()
## .. .. ..- attr(*, "class")= chr "collector_integer" "collector"
## .. ..$ name : list()
## .. .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ sex : list()
## .. .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ age : list()
## .. .. ..- attr(*, "class")= chr "collector_double" "collector"
## .. ..$ sibsp : list()
## .. .. ..- attr(*, "class")= chr "collector_integer" "collector"
## .. ..$ parch : list()
## .. .. ..- attr(*, "class")= chr "collector_integer" "collector"
## .. ..$ ticket : list()
## .. .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ fare : list()
## .. .. ..- attr(*, "class")= chr "collector_double" "collector"
## .. ..$ cabin : list()
## .. .. ..- attr(*, "class")= chr "collector_character" "collector"
## .. ..$ embarked: list()
## .. .. ..- attr(*, "class")= chr "collector_character" "collector"
## ..$ default: list()
## .. ..- attr(*, "class")= chr "collector_guess" "collector"
## ..- attr(*, "class")= chr "col_spec"
```

```

# Koristicemo samo kolone 'survived', 'pclass', 'sex' i 'age'
titanic <- titanic[, c(1, 2, 4, 5)]
str(titanic)

## Classes 'tbl_df', 'tbl' and 'data.frame':   891 obs. of  4 variables:
## $ survived: int  0 1 1 1 0 0 0 0 1 1 ...
## $ pclass : int  3 1 3 1 3 3 1 3 3 2 ...
## $ sex : chr "male" "female" "female" "female" ...
## $ age : num  22 38 26 35 35 NA 54 2 27 14 ...

# Prve tri promenljive bi evidentno trebalo da budu tretirane kao kategoricke promenljive - faktori
titanic[-4] <- map(titanic[-4], as.factor)

str(titanic)

## Classes 'tbl_df', 'tbl' and 'data.frame':   891 obs. of  4 variables:
## $ survived: Factor w/ 2 levels "0","1": 1 2 2 2 1 1 1 1 2 2 ...
## $ pclass : Factor w/ 3 levels "1","2","3": 3 1 3 1 3 3 1 3 3 2 ...
## $ sex : Factor w/ 2 levels "female","male": 2 1 1 1 2 2 2 2 1 1 ...
## $ age : num  22 38 26 35 35 NA 54 2 27 14 ...

table(titanic$survived)

##
## 0 1
## 549 342

# Odnos prezivelih i poginulih
prop.table(table(titanic$survived))

##
## 0 1
## 0.6161616 0.3838384

# Generisemo klasifikacioni model (drvo odlucivanja - decision tree) na osnovu datih podataka:
tree <- rpart(survived ~ ., data = titanic, method = "class")

# Koristimo predict() funkciju da predvidimo klase
pred <- predict(tree, newdata = titanic, type = "class")

# Konstruisemo konfuzionu matricu koristeći "table()":
conf_t <- table(titanic$survived, pred)
conf_t

## pred
## 0 1
## 0 479 70
## 1 94 248

```

## Primer 2:

```

#Isto to sa "pima" bazom podataka
library(faraway)

```

```

## Warning: package 'faraway' was built under R version 3.3.2
##

```

```
## Attaching package: 'faraway'

## The following object is masked from 'package:rpart':
##
##      solder

data(pima)

head(pima)

##   pregnant glucose diastolic triceps insulin  bmi diabetes age test
## 1         6     148         72      35       0 33.6    0.627  50    1
## 2         1      85         66      29       0 26.6    0.351  31    0
## 3         8     183         64       0       0 23.3    0.672  32    1
## 4         1      89         66      23      94 28.1    0.167  21    0
## 5         0     137         40      35     168 43.1    2.288  33    1
## 6         5     116         74       0       0 25.6    0.201  30    0

str(pima)

## 'data.frame':   768 obs. of  9 variables:
##  $ pregnant : int  6 1 8 1 0 5 3 10 2 8 ...
##  $ glucose   : int 148 85 183 89 137 116 78 115 197 125 ...
##  $ diastolic : int 72 66 64 66 40 74 50 0 70 96 ...
##  $ triceps   : int 35 29 0 23 35 0 32 0 45 0 ...
##  $ insulin   : int 0 0 0 94 168 0 88 0 543 0 ...
##  $ bmi       : num 33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
##  $ diabetes  : num 0.627 0.351 0.672 0.167 2.288 ...
##  $ age       : int 50 31 32 21 33 30 26 29 53 54 ...
##  $ test      : int 1 0 1 0 1 0 1 0 1 1 ...

# Da bismo obezbedili reproduktivnost
set.seed(33)

# Generisemo klasifikacioni model (drvo odlucivanja - decision tree) na osnovu datih podataka:
tree <- rpart(test ~ ., data = pima, method = "class")

# Koristimo predict() funkciju da predvidimo klase
pred <- predict(tree, newdata = pima, type = "class")

# Konstruisemo konfuzionu matricu koristeći "table()":
conf_p <- table(pima$test, pred)
conf_p

##      pred
##      0    1
## 0 449   51
## 1   72 196
```

### Tacnost, preciznost, senzitivnost (recall), specificnost - Primer

```
# Izracunajmo parametre za ocenu valjanosti modela "tree" za "titanic" skup podataka

# Formiramo TP, FN, FP i TN na osnovu "conf_t"

TP <- conf_t[2,2]
```

```

FP <- conf_t[1,2]

FN <- conf_t[2,1]

TN <- conf_t[1,1]

# Tacnost (Accuracy)
acc <- (TP + TN)/sum(conf_t)
acc

## [1] 0.8159371

# Preciznost (Precision)
prec <- TP/(TP + FP)
prec

## [1] 0.7798742

# Senzitivnost (Sensitivity, Recall)
sens <- TP/(TP + FN)
sens

## [1] 0.7251462

# Specifichost (Specificity)
spec <- TN/(TN + FP)
spec

## [1] 0.8724954

```

### Zadatak za vezbanje na casu:

Izracunajte ove vrednosti za “tree” model generisan na osnovu “pima” seta podataka.

### Kvalitet regresije

- Srednja kvadratna greska
- U nasem slucaju mozemo smatrati da se poklapa sa standardnom devijacijom
- $\sqrt{(1/nrow(truth)) * \sum( (truth\$col - pred)^2)}$

### Primer:

```

# Koristicemo "pima" bazu

# Struktura seta podataka
str(pima)

## 'data.frame':  768 obs. of  9 variables:
## $ pregnant : int  6 1 8 1 0 5 3 10 2 8 ...
## $ glucose  : int  148 85 183 89 137 116 78 115 197 125 ...
## $ diastolic: int  72 66 64 66 40 74 50 0 70 96 ...
## $ triceps  : int  35 29 0 23 35 0 32 0 45 0 ...
## $ insulin  : int  0 0 0 94 168 0 88 0 543 0 ...
## $ bmi      : num  33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...

```

```
## $ diabetes : num  0.627 0.351 0.672 0.167 2.288 ...
## $ age      : int  50 31 32 21 33 30 26 29 53 54 ...
## $ test     : int   1 0 1 0 1 0 1 0 1 1 ...

# Multivarijabilna linearna regresija - prostiji model (uključen manji broj promenljivih)
fit_1 <- lm(diabetes ~ bmi + triceps + age + glucose, data = pima)

# Predviđanje na osnovu modela: pred_1
pred_1 <- predict(fit_1)

# RMSE na osnovu "pima$diabetes" (tačne vrednosti) i "pred_1" (vrednosti na osnovu modela fit_1)
rmse_1 <- sqrt(1/nrow(pima)*sum((pima$diabetes - pred_1) ^ 2))

rmse_1

## [1] 0.3222776

# Multivarijabilna linearna regresija - kompleksniji model (uključen veći broj promenljivih)
fit_2 <- lm(diabetes ~ bmi + triceps + age + glucose + diastolic + insulin + pregnant, data = pima)

# Predviđanje na osnovu modela: pred_1
pred_2 <- predict(fit_2)

# RMSE na osnovu "pima$diabetes" (tačne vrednosti) i "pred_1" (vrednosti na osnovu modela fit_1)
rmse_2 <- sqrt(1/nrow(pima)*sum((pima$diabetes - pred_2) ^ 2))

rmse_2

## [1] 0.3205351
```

## Procena valjanosti klasterizacije: WSS vs BSS

```
# Da bi smo obezbedili reproduktivnost
set.seed(33)

# Proveravamo strukturu podataka
str(iris)

## 'data.frame':  150 obs. of  5 variables:
## $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

head(iris)

##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1         5.1         3.5         1.4         0.2  setosa
## 2         4.9         3.0         1.4         0.2  setosa
## 3         4.7         3.2         1.3         0.2  setosa
## 4         4.6         3.1         1.5         0.2  setosa
## 5         5.0         3.6         1.4         0.2  setosa
## 6         5.4         3.9         1.7         0.4  setosa
```



```
# Delimo "iris" na dva seta: "my_iris" i "species"
my_iris <- iris[-5]
species <- iris$Species

# Vrsimo k-means klasterizaciju za "my_iris" uz pretpostavku da postoje tri klase: "kmeans_iris"
kmeans_iris <- kmeans(my_iris,3)

# Poredimo dobijene klasterne sa istinskim klasama (kategorijama)
table(species, kmeans_iris$cluster)

##
## species      1  2  3
## setosa       0  0 50
## versicolor  48  2  0
## virginica   14 36  0

# Plotujemo "Petal.Width" vs "Petal.Length", bojimo po klasterima odn. postojećim kategorijama
par(mfrow = c(1,2))
plot(Petal.Length ~ Petal.Width, data = my_iris, col = kmeans_iris$cluster)
title("k-means - klasteri")
plot(Petal.Length ~ Petal.Width, data = my_iris, col = iris$Species)
title("Istinske klase")
```



```
kmeans_iris$tot.withinss/kmeans_iris$betweenss
```

```
## [1] 0.1308696
```

## Trening set i test set

- Cilj implementacije algoritma **nadgledanog** učenja jeste dobijanje “dovoljno” dobrog prediktivnog modela na osnovu raspoloživog seta podataka.
- Set podataka koji se koristi za formiranje modela - **trening set**
- Set podatak koji se koristi za procenu valjanosti modela - **test set**
- Trening set i test set ne smeju imati/deliti zajednicke elemente tj. opservacije
- Samo testiranjem modela na podacima koji nisu korisceni za učenje mozemo izvesti adekvatnu estimaciju ocena valjanosti modela - generalizacija.
- Opste prihvacena praksa je da se rasplziv skup podataka podeli na sledeci nacin:
  - Trening set 70% ili 75%
  - Test set 30% ili 35%
- Prilikom podele raspoloživog skupa podataka treba strogo voditi racuna da zastupljenost, odn. distribucija, klasa (ovo se odnosi na algoritme za klasifikaciju) bude slicna u trening i test setu
  - ne bi smelo da se dogodi da jedan ili drugi set uopste ne sadrže ni jednu opservaciju koja pripada određenoj klasi
- Dobra praksa je da se poredak opservacija randomizuje (slučajno odabrana permutacija) pre deljenja skupa podataka na trening i test set
  - Ovo vazi i za klasifikaciju i za regresiju
- Odabiranje (semplovanje) opservacija za trening i test set može ponekad i značajno uticati na procenjene vrednosti ocena valjanosti datog modela
  - Da bi se ovaj efekat minimizovao koristi se **unakrsna validacija** (cross-validation)

## Primer

```
# Koristicemo "titanic" set podataka formiran u jednom od prethodnih primera
str(titanic)

## Classes 'tbl_df', 'tbl' and 'data.frame':   891 obs. of  4 variables:
## $ survived: Factor w/ 2 levels "0","1": 1 2 2 2 1 1 1 1 2 2 ...
## $ pclass  : Factor w/ 3 levels "1","2","3": 3 1 3 1 3 3 1 3 3 2 ...
## $ sex      : Factor w/ 2 levels "female","male": 2 1 1 1 2 2 2 2 1 1 ...
## $ age      : num  22 38 26 35 35 NA 54 2 27 14 ...

table(titanic$survived)

##
##    0    1
## 549 342

# Odnos prezivelih i poginulih
prop.table(table(titanic$survived))

##
##          0          1
## 0.6161616 0.3838384

# Da bismo omogucili reproduktivnost
set.seed(33)

# Prvo napravimo jednu slucajno odabranu permutaciju celog skupa podataka (dataset shuffle)
n <- nrow(titanic)
shuffled <- titanic[sample(n),] #f-a 'sample' vrsi slucajno odabiranje elemenata zadanog vektora

# Delimo skup podataka na trening i test set (70% i 30%)
```

```

train_indicies <- 1:round(0.7 * n)
train <- shuffled[train_indicies, ]
test <- shuffled[-train_indicies, ]

# Generisemo klasifikacioni model (drvo odlucivanja - decision tree) na osnovu trening seta:
tree <- rpart(survived ~ ., data = train, method = "class")

# Koristeci dobijeni model "tree" vrsimo klasifikaciju podataka iz test seta:
pred <- predict(tree, newdata = test, type = "class")

# Racunamo matricu konfuzije
conf_t <- table(test$survived, pred)

# Prikaz matrice konfuzije
conf_t

##      pred
##      0   1
## 0 128  28
## 1  36  75

# Formiramo TP, FN, FP i TN na osnovu "conf_t"

TP <- conf_t[2,2]

FP <- conf_t[1,2]

FN <- conf_t[2,1]

TN <- conf_t[1,1]

# Tacnost (Accuracy)
acc <- (TP + TN)/sum(conf_t)
acc

## [1] 0.7602996

# Preciznost (Precision)
prec <- TP/(TP + FP)
prec

## [1] 0.7281553

# Senzitivnost (Sensitivity, Recall)
sens <- TP/(TP + FN)
sens

## [1] 0.6756757

# Specifcnost (Specificity)
spec <- TN/(TN + FP)
spec

## [1] 0.8205128

```



**Zadatak za vezbanje na casu:**

Ponovite pokazanu proceduru koristeći "pima" skup podataka.

### Upotreba unakrsne validacije (cross-validation)

Radi demonstracije ćemo ručno formirati algoritam koji koristi unakrsnu validaciju za procenu tačnosti modela:

```
# Da bismo obezbedili reproduktivnost
set.seed(33)

# Koristicemo prethodno formirani "shuffled" skup podataka

# Inicijalizujemo vektor accs - popunjavamo nulama
accs <- rep(0,9)

# Treniramo model koristeći kros-validacione intervale vrednosti i vrsimo estimaciju tacnosti modela ka
for (i in 1:9) {
  # Ovi indeksi ukazuju na trenutni interval test seta koji koristimo za treniranje modela
  indices <- (((i - 1) * round((1/9)*nrow(shuffled))) + 1):((i*round((1/9) * nrow(shuffled))))

  # Iskljuccujemo ove intervale iz trening seta
  train <- shuffled[-indices,]

  # Ukljucimo ih u test set
  test <- shuffled[indices,]

  # Treniramo model sa svakim od dobijenih trening setova po iteracijama
  tree <- rpart(survived ~ ., train, method = "class")

  # Predvidjamo klase za tekuci test set u svakoj od iteracija
  pred <- predict(tree, test, type = "class")

  # Formiramo odgovarajucu konfuzionu matricu
  conf <- table(test$survived, pred)

  # Dodeljujemo vrednost za tacnost tekuceg modela i-tom indeksu u vektoru accs
  accs[i] <- sum(diag(conf))/sum(conf)
}

# Srednja vrednost za accs
mean(accs)
```

```
## [1] 0.7833895
```

**Pitanje:** Recimo da primenjujemo unakrsnu validaciju na skupu podataka koji sadrži 22680 opservacija. Zelite da vas trening set sadrži 21420 unosa (opservacija). Koliko iteracija može da sadrži kros-validacioni algoritam?

### Bajas i varijansa (Bias and Variance)

#### Primer

Koristicemo *Spambase Data Set* koji mozete naci na <https://archive.ics.uci.edu/ml/datasets/Spambase>

```
if (!"emails_full" %in% ls()) {  
  emails_full <- read.csv("data/spambase.data", header = FALSE)  
}
```

```
# Proveravamo strukturu seta podataka  
str(emails_full)
```

```
## 'data.frame': 4601 obs. of 58 variables:  
## $ V1 : num 0 0.21 0.06 0 0 0 0 0 0.15 0.06 ...  
## $ V2 : num 0.64 0.28 0 0 0 0 0 0 0.12 ...  
## $ V3 : num 0.64 0.5 0.71 0 0 0 0 0 0.46 0.77 ...  
## $ V4 : num 0 0 0 0 0 0 0 0 0 ...  
## $ V5 : num 0.32 0.14 1.23 0.63 0.63 1.85 1.92 1.88 0.61 0.19 ...  
## $ V6 : num 0 0.28 0.19 0 0 0 0 0 0.32 ...  
## $ V7 : num 0 0.21 0.19 0.31 0.31 0 0 0 0.3 0.38 ...  
## $ V8 : num 0 0.07 0.12 0.63 0.63 1.85 0 1.88 0 0 ...  
## $ V9 : num 0 0 0.64 0.31 0.31 0 0 0 0.92 0.06 ...  
## $ V10: num 0 0.94 0.25 0.63 0.63 0 0.64 0 0.76 0 ...  
## $ V11: num 0 0.21 0.38 0.31 0.31 0 0.96 0 0.76 0 ...  
## $ V12: num 0.64 0.79 0.45 0.31 0.31 0 1.28 0 0.92 0.64 ...  
## $ V13: num 0 0.65 0.12 0.31 0.31 0 0 0 0.25 ...  
## $ V14: num 0 0.21 0 0 0 0 0 0 0 ...  
## $ V15: num 0 0.14 1.75 0 0 0 0 0 0.12 ...  
## $ V16: num 0.32 0.14 0.06 0.31 0.31 0 0.96 0 0 0 ...  
## $ V17: num 0 0.07 0.06 0 0 0 0 0 0 ...  
## $ V18: num 1.29 0.28 1.03 0 0 0 0.32 0 0.15 0.12 ...  
## $ V19: num 1.93 3.47 1.36 3.18 3.18 0 3.85 0 1.23 1.67 ...  
## $ V20: num 0 0 0.32 0 0 0 0 0 3.53 0.06 ...  
## $ V21: num 0.96 1.59 0.51 0.31 0.31 0 0.64 0 2 0.71 ...  
## $ V22: num 0 0 0 0 0 0 0 0 0 ...  
## $ V23: num 0 0.43 1.16 0 0 0 0 0 0.19 ...  
## $ V24: num 0 0.43 0.06 0 0 0 0 0 0.15 0 ...  
## $ V25: num 0 0 0 0 0 0 0 0 0 ...  
## $ V26: num 0 0 0 0 0 0 0 0 0 ...  
## $ V27: num 0 0 0 0 0 0 0 0 0 ...  
## $ V28: num 0 0 0 0 0 0 0 0 0 ...  
## $ V29: num 0 0 0 0 0 0 0 0 0 ...  
## $ V30: num 0 0 0 0 0 0 0 0 0 ...  
## $ V31: num 0 0 0 0 0 0 0 0 0 ...  
## $ V32: num 0 0 0 0 0 0 0 0 0 ...  
## $ V33: num 0 0 0 0 0 0 0 0 0.15 0 ...  
## $ V34: num 0 0 0 0 0 0 0 0 0 ...  
## $ V35: num 0 0 0 0 0 0 0 0 0 ...  
## $ V36: num 0 0 0 0 0 0 0 0 0 ...  
## $ V37: num 0 0.07 0 0 0 0 0 0 0 ...  
## $ V38: num 0 0 0 0 0 0 0 0 0 ...  
## $ V39: num 0 0 0 0 0 0 0 0 0 ...  
## $ V40: num 0 0 0.06 0 0 0 0 0 0 ...  
## $ V41: num 0 0 0 0 0 0 0 0 0 ...  
## $ V42: num 0 0 0 0 0 0 0 0 0 ...  
## $ V43: num 0 0 0.12 0 0 0 0 0 0.3 0 ...  
## $ V44: num 0 0 0 0 0 0 0 0 0.06 ...  
## $ V45: num 0 0 0.06 0 0 0 0 0 0 ...
```

```
## $ V46: num 0 0 0.06 0 0 0 0 0 0 0 ...
## $ V47: num 0 0 0 0 0 0 0 0 0 0 ...
## $ V48: num 0 0 0 0 0 0 0 0 0 0 ...
## $ V49: num 0 0 0.01 0 0 0 0 0 0 0.04 ...
## $ V50: num 0 0.132 0.143 0.137 0.135 0.223 0.054 0.206 0.271 0.03 ...
## $ V51: num 0 0 0 0 0 0 0 0 0 0 ...
## $ V52: num 0.778 0.372 0.276 0.137 0.135 0 0.164 0 0.181 0.244 ...
## $ V53: num 0 0.18 0.184 0 0 0 0.054 0 0.203 0.081 ...
## $ V54: num 0 0.048 0.01 0 0 0 0 0 0.022 0 ...
## $ V55: num 3.76 5.11 9.82 3.54 3.54 ...
## $ V56: int 61 101 485 40 40 15 4 11 445 43 ...
## $ V57: int 278 1028 2259 191 191 54 112 49 1257 749 ...
## $ V58: int 1 1 1 1 1 1 1 1 1 1 ...
```

*# Na osnovu dokumentacije...*

```
emails_full <- emails_full[, c(55, 58)]
```

```
str(emails_full)
```

```
## 'data.frame': 4601 obs. of 2 variables:
## $ V55: num 3.76 5.11 9.82 3.54 3.54 ...
## $ V58: int 1 1 1 1 1 1 1 1 1 1 ...
```

```
colnames(emails_full) <- c("avg_capital_seq", "spam")
```

```
str(emails_full)
```

```
## 'data.frame': 4601 obs. of 2 variables:
## $ avg_capital_seq: num 3.76 5.11 9.82 3.54 3.54 ...
## $ spam : int 1 1 1 1 1 1 1 1 1 1 ...
```

```
emails_full$spam <- as.factor(emails_full$spam)
```

```
str(emails_full)
```

```
## 'data.frame': 4601 obs. of 2 variables:
## $ avg_capital_seq: num 3.76 5.11 9.82 3.54 3.54 ...
## $ spam : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
```

*# Definisemo funkciju spam\_classifier()*

*# 1 - spam, 0 - ham*

```
spam_classifier <- function(x){
  prediction <- rep(NA,length(x))
  prediction[x > 4] <- 1
  prediction[x >= 3 & x <= 4] <- 0
  prediction[x >= 2.2 & x < 3] <- 1
  prediction[x >= 1.4 & x < 2.2] <- 0
  prediction[x > 1.25 & x < 1.4] <- 1
  prediction[x <= 1.25] <- 0
  return(factor(prediction, levels = c("0","1")))
}
```

*# Primenimo spam\_classifier na emails\_full: pred\_full*

```
pred_full <- spam_classifier(emails_full$avg_capital_seq)
```

*# Konfuzionarna matrica za emails\_full: conf\_full*

```

conf_full <- table(emails_full$spam, pred_full)

# Racunamo tacnost na osnovu conf_full: acc_full
acc_full <- sum(diag(conf_full))/sum(conf_full)
acc_full

## [1] 0.6561617

# Uproscen model za klasifikaciju
spam_classifier <- function(x){
  prediction <- rep(NA,length(x))
  prediction[x > 4] <- 1
  prediction[x <= 4] <- 0
  return(factor(prediction, levels = c("0","1")))
}

# Tacnost predikcije sa uproscenim modelom za emails data set
conf_small <- table(emails$spam, spam_classifier(emails$avg_capital_seq))
acc_small <- sum(diag(conf_small)) / sum(conf_small)
acc_small

## [1] 0.7692308

# Primenimo uprosceni model i na "emails_full" i sracunamo konfuzionu matricu
conf_full <- table(emails_full$spam, spam_classifier(emails_full$avg_capital_seq))

# Izracunamo tacnost
acc_full <- sum(diag(conf_full)) / sum(conf_full)
acc_full

## [1] 0.7259291

```

## Regresija

### Prosta linearna regresija

#### Koriscenje funkcije lm() - Primeri

##### Primer 1:

U ovom primeru cemo koristiti set podataka "Boston" iz paketa MASS. Ovaj set podataka sadrzi podatke o trzisnoj vrednosti nekretnina u predgradjima Bostona, SAD, zajedno sa razlicitim parametrima koji uticu na formiranje ove vrednosti.

```

library(MASS)

##
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##      select

library(ISLR)
library(ggplot2)

?Boston

```

```
## starting httpd help server ...
```

```
## done
```

```
str(Boston)
```

```
## 'data.frame':  506 obs. of  14 variables:
## $ crim   : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn     : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus  : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas   : int   0 0 0 0 0 0 0 0 0 0 ...
## $ nox    : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm     : num  6.58 6.42 7.18 7 7.15 ...
## $ age    : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis    : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad    : int   1 2 2 3 3 3 5 5 5 ...
## $ tax    : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black  : num  397 397 393 395 397 ...
## $ lstat  : num  4.98 9.14 4.03 2.94 5.33 ...
## $ medv   : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

```
head(Boston)
```

```
##      crim zn indus chas   nox    rm age    dis rad tax ptratio black
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900   1 296    15.3 396.90
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671   2 242    17.8 396.90
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671   2 242    17.8 392.83
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622   3 222    18.7 394.63
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622   3 222    18.7 396.90
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622   3 222    18.7 394.12
##   lstat medv
## 1  4.98 24.0
## 2  9.14 21.6
## 3  4.03 34.7
## 4  2.94 33.4
## 5  5.33 36.2
## 6  5.21 28.7
```

```
# Proverimo kako izgleda promena "medv" (median value of owner-occupied homes in \($1000s) sa
# "lstat" (lower status of the population (percent))
```

```
plot(medv~lstat,Boston)
```

```
# Kao sto vidimo postoji jasan trend opadanja vrednosti nekretnina sa porastom procenta
# siromasnijih stanovnika (ovakva korelacija je naravno i ocekivana). Ovakvi slucajevi su dobri
# kandidati za modelovanje prostom linerarnom regresijom.
```

```
fit_1 = lm(medv ~ lstat, data = Boston)
```

```
# Hajde da vidimo kako izgleda nas model
fit_1
```

```
##
```

```
## Call:
```

```
## lm(formula = medv ~ lstat, data = Boston)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)      lstat
##      34.55      -0.95

# Detaljniji uvid
summary(fit_1)

##
## Call:
## lm(formula = medv ~ lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.168  -3.990  -1.318   2.034  24.500
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 34.55384    0.56263   61.41  <2e-16 ***
## lstat      -0.95005    0.03873  -24.53  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.216 on 504 degrees of freedom
## Multiple R-squared:  0.5441, Adjusted R-squared:  0.5432
## F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16


# Sta sve model sadrzi
names(fit_1)

## [1] "coefficients" "residuals"      "effects"        "rank"
## [5] "fitted.values" "assign"          "qr"             "df.residual"
## [9] "xlevels"      "call"           "terms"          "model"

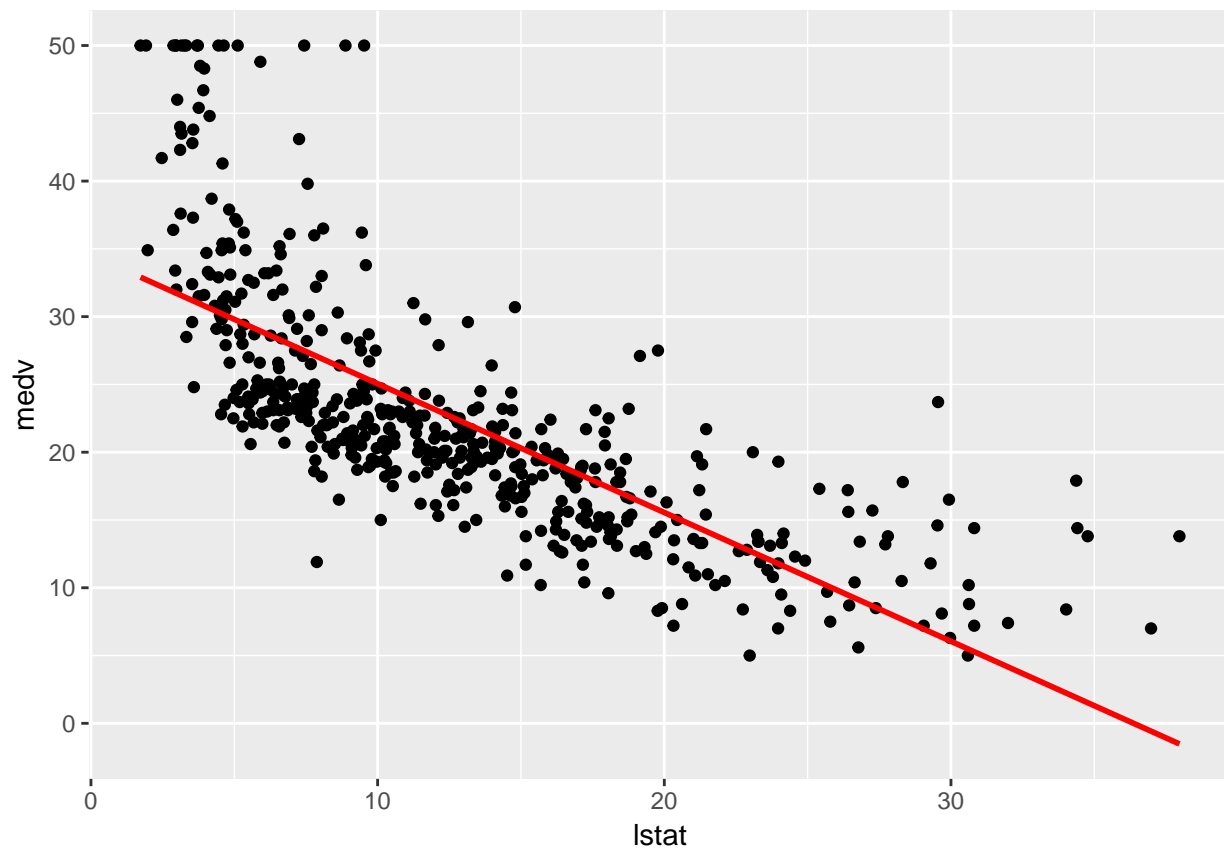
# Samo koeficijenti
fit_1$coefficients

## (Intercept)      lstat
## 34.5538409  -0.9500494

# Ucrtažmo regresionu pravu na početni scatter plot
abline(fit_1$coefficients, col = "red", lwd = 2)
```

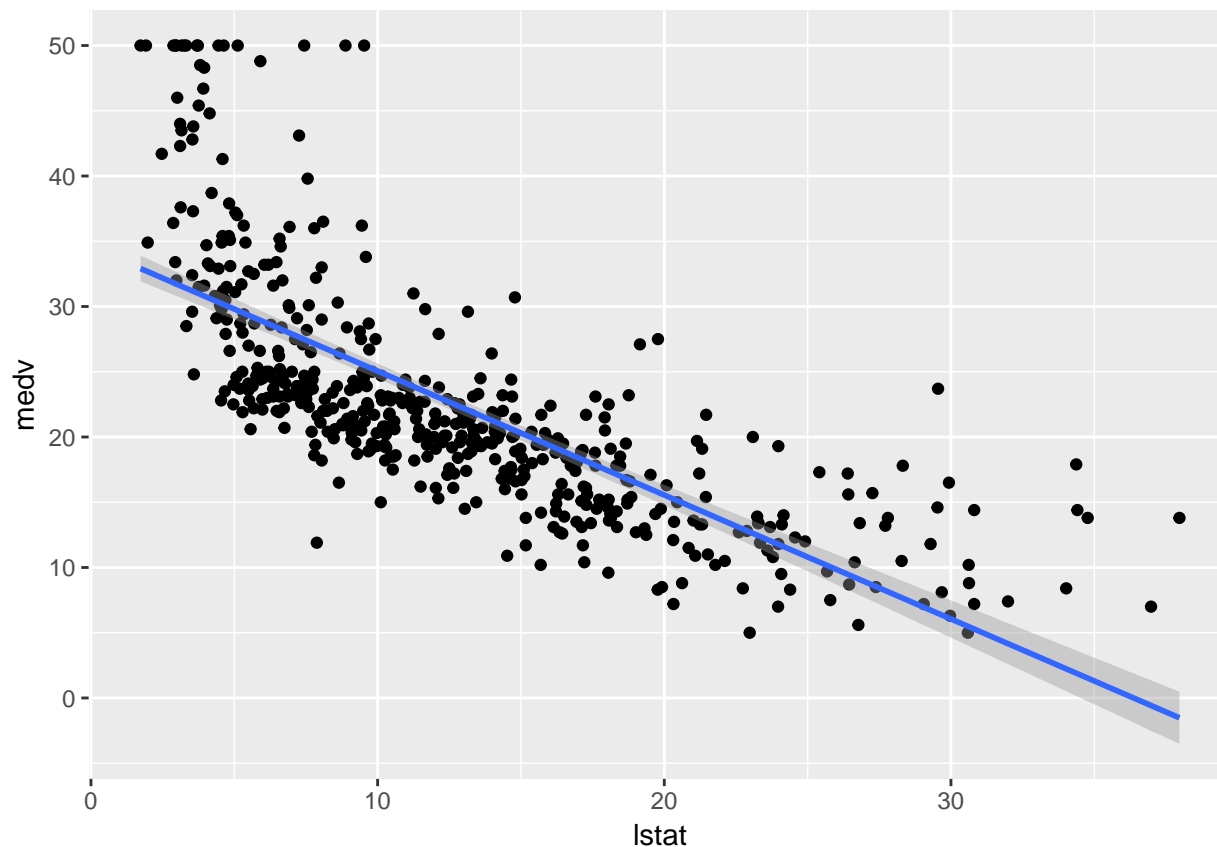


```
# Ili sve zajedno koristerci "ggplot2"  
ggplot(Boston, aes( x = lstat, y = medv)) +  
  geom_point() +  
  geom_smooth(method = "lm", se = FALSE, colour = "red")
```



```
# Ako zelimo i interval pouzdanosti sam izostavimo parametar "se" (podrazumevano se = TRUE)  
ggplot(Boston, aes( x = lstat, y = medv)) +  
  geom_point() +  
  geom_smooth(method = "lm")
```





```
# Interval pouzdanosti
confint(fit_1)
```

```
##              2.5 %      97.5 %
## (Intercept) 33.448457 35.6592247
## lstat      -1.026148 -0.8739505
```

```
# Da prevedimo vrednosti "medv" za dati vektor vrednosti "lstat" promenljive, uz
# proračun intervala pouzdanosti
predict(fit_1,data.frame(lstat = c(5,10,15)),interval = "confidence")
```

```
##      fit      lwr      upr
## 1 29.80359 29.00741 30.59978
## 2 25.05335 24.47413 25.63256
## 3 20.30310 19.73159 20.87461
```

## Multivarijabilna linearna regresija

### Primer 1

```
library(readr)
library(tidyr)
library(purrr)
library(ggpubr)
```

```
## Warning: package 'ggpubr' was built under R version 3.3.2
```

```

# Uvoz i sredjivanje podataka
shop_data <- read_csv("data/shop_data.csv")

## Parsed with column specification:
## cols(
##   `sales`,`sq_ft`,`inv`,`ads`,`size_dist`,`comp` = col_character()
## )

shop_data <- separate(shop_data, "sales","sq_ft","inv","ads","size_dist","comp",
                       c("sales","sq_ft","inv","ads","size_dist","comp"), sep = ",")
shop_data <- as.data.frame(map(shop_data, as.numeric))

str(shop_data)

## 'data.frame':   27 obs. of  6 variables:
## $ sales      : num  231 156 10 519 437 487 299 195 20 68 ...
## $ sq_ft      : num   3 2.2 0.5 5.5 4.4 ...
## $ inv        : num  294 232 149 600 567 571 512 347 212 102 ...
## $ ads        : num   8.2 6.9 3 12 10.6 ...
## $ size_dist  : num   8.2 4.1 4.3 16.1 14.1 ...
## $ comp       : num   11 12 15 1 5 4 10 12 15 8 ...

head(shop_data)

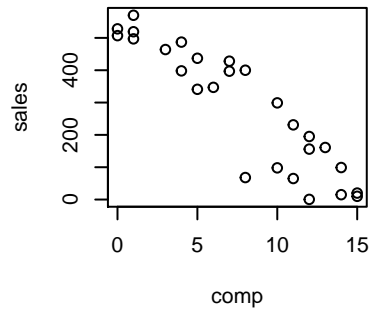
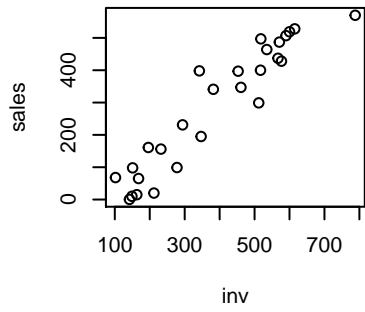
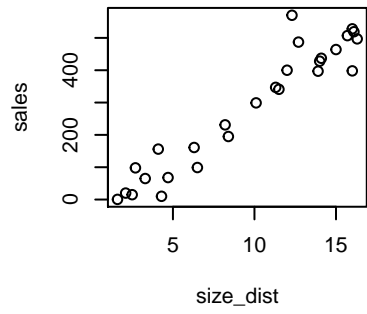
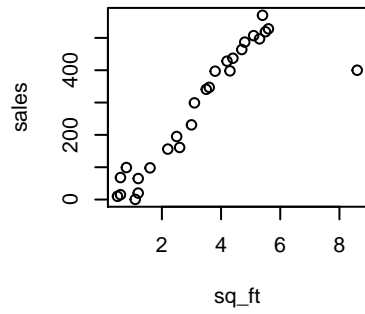
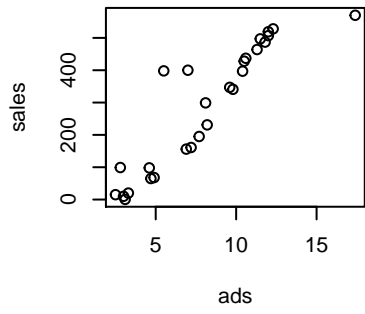
##   sales sq_ft inv  ads size_dist comp
## 1   231   3.0 294  8.2         8.2  11
## 2   156   2.2 232  6.9         4.1  12
## 3    10   0.5 149  3.0         4.3  15
## 4   519   5.5 600 12.0        16.1   1
## 5   437   4.4 567 10.6        14.1   5
## 6   487   4.8 571 11.8        12.7   4

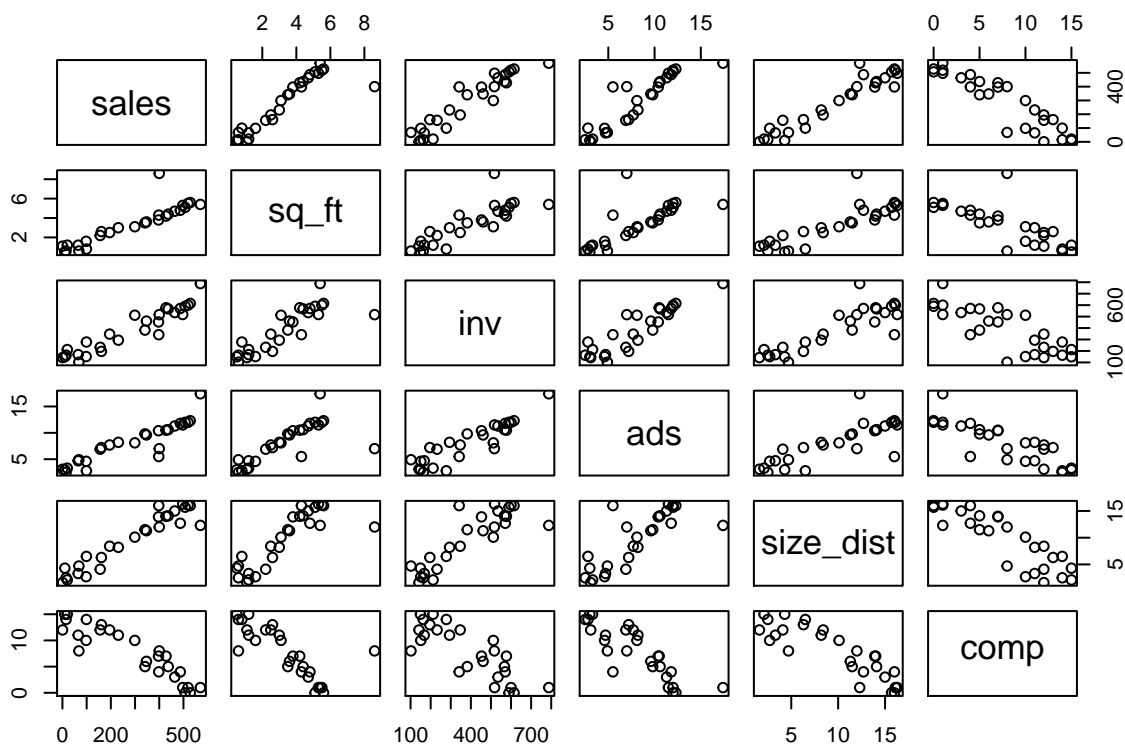
# Hajde da proverimo kako se podaci ponasaju i mogu li se uociti relacije
# izmedju distribucija promenljivih koje bi ukazivale na opravdanost uvodjenja
# linearnog modela:

par(mfrow = c(2,3))
plot(sales ~ ads, shop_data)
plot(sales ~ sq_ft, shop_data)
plot(sales ~ size_dist, shop_data)
plot(sales ~ inv, shop_data)
plot(sales ~ comp, shop_data)

#Ili:
pairs(shop_data)

```





```
# Linearni model za "sales" koji uključuje sve prediktore (sve preostale promenljive)
lm_shop_1 <- lm( sales ~., data = shop_data)

# Proverimo parametre valjanosti modela i koliko su pojedini prediktori znacajni u modelu:
summary(lm_shop_1)
```

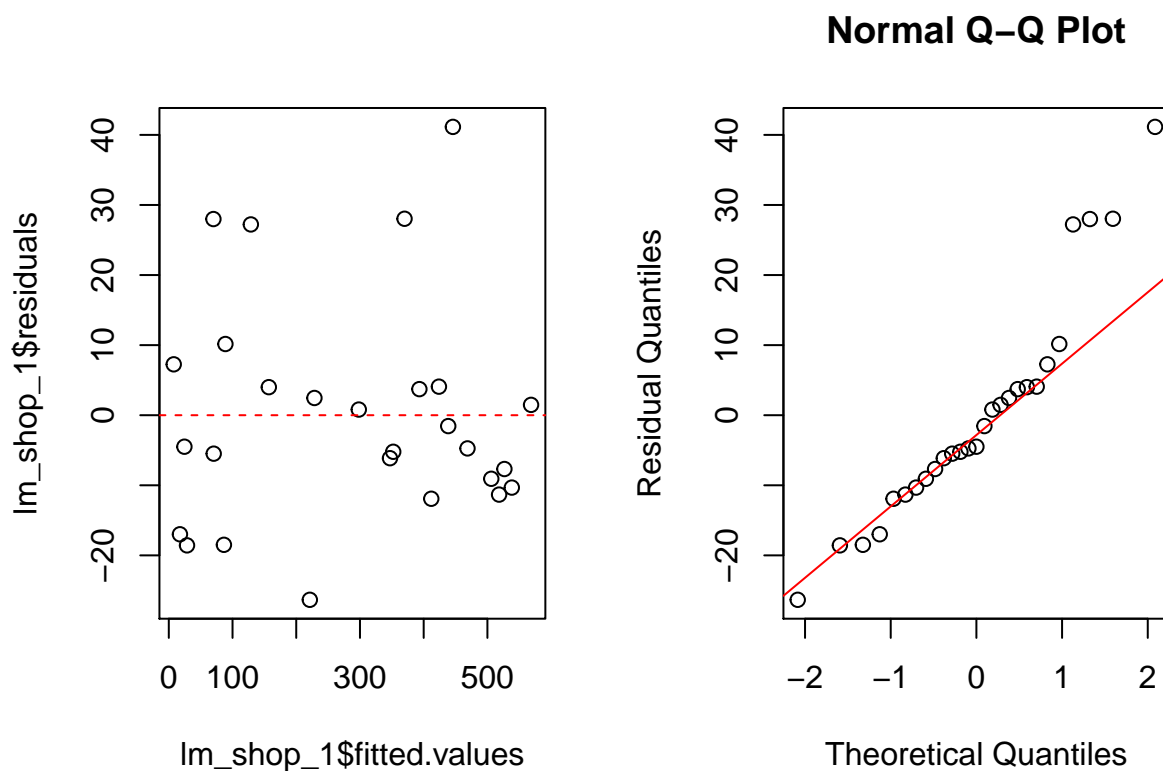
```
##
## Call:
## lm(formula = sales ~ ., data = shop_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -26.338  -9.699  -4.496   4.040  41.139
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -18.85941   30.15023  -0.626  0.538372
## sq_ft       16.20157    3.54444   4.571  0.000166 ***
## inv         0.17464    0.05761   3.032  0.006347 **
## ads         11.52627    2.53210   4.552  0.000174 ***
## size_dist   13.58031    1.77046   7.671  1.61e-07 ***
## comp        -5.31097    1.70543  -3.114  0.005249 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 17.65 on 21 degrees of freedom
```

```
## Multiple R-squared:  0.9932, Adjusted R-squared:  0.9916
## F-statistic: 611.6 on 5 and 21 DF,  p-value: < 2.2e-16
```

*# Da bismo uopste mogli da koristimo p-vrednosti u ovom kontekstu treba prvo da proverimo  
# da li je zadovoljena pretpostavka o normalnoj distribuciji reziduala!*

```
par(mfrow = c(1,2))
# Plotujemo rezidualne u funkciji fitovanih vrednosti za pojedinačne opservacije
plot(lm_shop_1$fitted.values, lm_shop_1$residuals)
abline(0,0, col = "red", lty = 2)


# Napravimo Q-Q plot kvantila reziduala
qqnorm(lm_shop_1$residuals, ylab = "Residual Quantiles")
qqline(lm_shop_1$residuals, col = "red")
```



```
par(mfrow = c(1,1))

# Mozemo i da upotrebimo f-ju "ggqqplot" iz paketa "ggpubr" koji sadrzi funkcije za
# plotovanje "lepih" grafika:

ggqqplot(lm_shop_1$residuals, ylab = "Residual Quantiles")
```



*# Me može se uočiti nikakav jasan "pattern" u distribucij reziduala, šta više kvantili  
# reziduala su uglavnom na liniji koja odgovara teorijskoj - normalnoj distribuciji*

*# Proverimo ponovo summary*

```
summary(lm_shop_1)
```

```
##
## Call:
## lm(formula = sales ~ ., data = shop_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -26.338  -9.699  -4.496   4.040  41.139
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -18.85941   30.15023  -0.626  0.538372
## sq_ft        16.20157    3.54444   4.571  0.000166 ***
## inv          0.17464    0.05761   3.032  0.006347 **
## ads          11.52627    2.53210   4.552  0.000174 ***
## size_dist    13.58031    1.77046   7.671  1.61e-07 ***
## comp         -5.31097    1.70543  -3.114  0.005249 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 17.65 on 21 degrees of freedom
```

```
## Multiple R-squared:  0.9932, Adjusted R-squared:  0.9916
## F-statistic: 611.6 on 5 and 21 DF,  p-value: < 2.2e-16
# Iskoristimo sada dobijeni model da predvidimo neto prodajnu vrednost na osnovu novog
# skupa prediktora:
shop_new = data.frame("sq_ft" = 2.3, "inv" = 420, "ads" = 8.7,
                      "size_dist" = 9.1, "comp" = 10)
predict(lm_shop_1, shop_new)

##          1
## 262.5006
```

## Primer 2

Za ovaj primer cemo ponovo koristiti set podataka "Boston" iz paketa MASS.

```
# Linearni model za "medv" na osnovu dva prediktora: "lstat" i "age"
fit2 = lm(medv ~ lstat + age, data = Boston)
summary(fit2)
```

```
##
## Call:
## lm(formula = medv ~ lstat + age, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.981  -3.978  -1.283   1.968  23.158
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  33.22276    0.73085  45.458 < 2e-16 ***
## lstat       -1.03207    0.04819 -21.416 < 2e-16 ***
## age          0.03454    0.01223   2.826  0.00491 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.173 on 503 degrees of freedom
## Multiple R-squared:  0.5513, Adjusted R-squared:  0.5495
## F-statistic:   309 on 2 and 503 DF,  p-value: < 2.2e-16
```

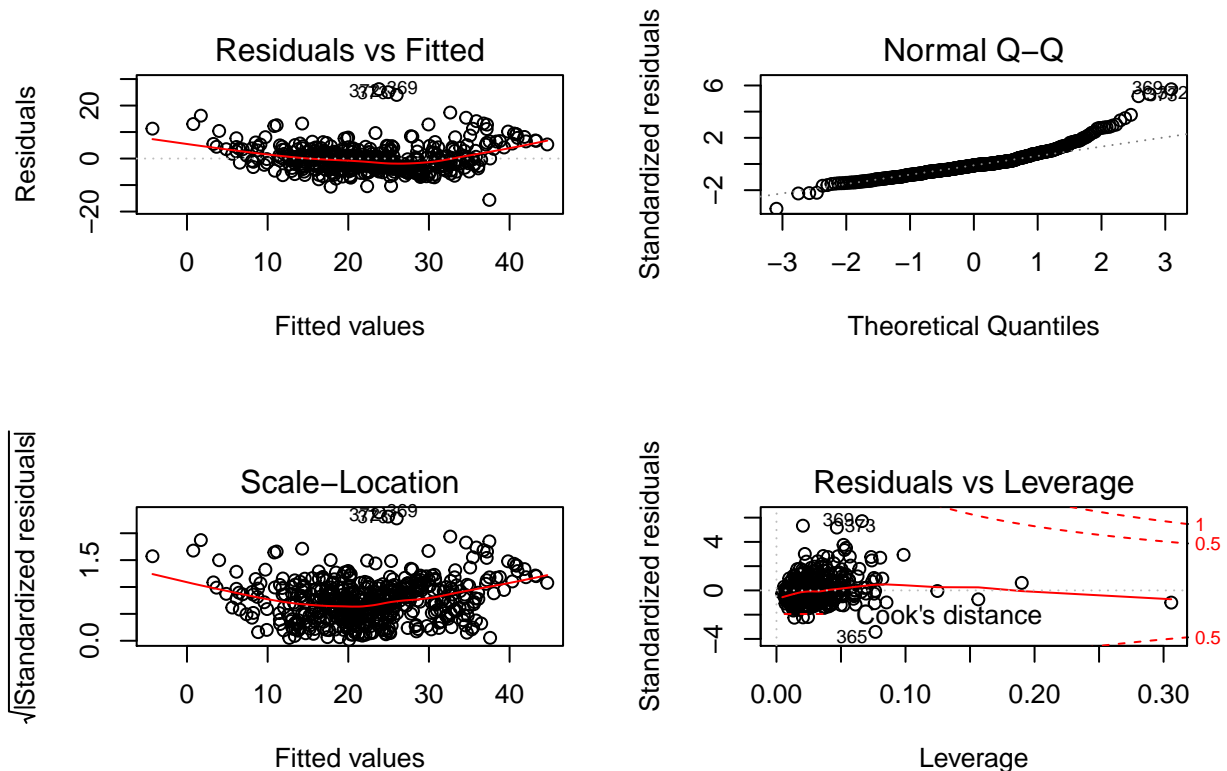
```
#Linearni model za "medv" na osnovu svih raspolozivih prediktora
fit3 = lm(medv ~.,Boston)
summary(fit3)
```

```
##
## Call:
## lm(formula = medv ~ ., data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.595  -2.730  -0.518   1.777  26.199
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.646e+01  5.103e+00   7.144 3.28e-12 ***
## crim        -1.080e-01  3.286e-02  -3.287 0.001087 **
```

```
## zn          4.642e-02  1.373e-02   3.382 0.000778 ***
## indus       2.056e-02  6.150e-02   0.334 0.738288
## chas        2.687e+00  8.616e-01   3.118 0.001925 **
## nox        -1.777e+01  3.820e+00  -4.651 4.25e-06 ***
## rm          3.810e+00  4.179e-01   9.116 < 2e-16 ***
## age         6.922e-04  1.321e-02   0.052 0.958229
## dis        -1.476e+00  1.995e-01  -7.398 6.01e-13 ***
## rad         3.060e-01  6.635e-02   4.613 5.07e-06 ***
## tax        -1.233e-02  3.760e-03  -3.280 0.001112 **
## ptratio     -9.527e-01  1.308e-01  -7.283 1.31e-12 ***
## black       9.312e-03  2.686e-03   3.467 0.000573 ***
## lstat      -5.248e-01  5.072e-02 -10.347 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.745 on 492 degrees of freedom
## Multiple R-squared:  0.7406, Adjusted R-squared:  0.7338
## F-statistic: 108.1 on 13 and 492 DF,  p-value: < 2.2e-16
```

*# Jos jedan nacín da se iscraju grafici koji se koriste za procenu valjanosti i opravdanosti linearnog modela*

```
par(mfrow = c(2,2))
plot(fit3)
```



*# Na osnovu "summary" za model fit3 videli smo da promenljive "indus" i "age" ne igraju bitnu ulogu, te*

```
fit4 = update(fit3, ~. - age - indus)
summary(fit4)
```



```
##
## Call:
## lm(formula = medv ~ crim + zn + chas + nox + rm + dis + rad +
##      tax + ptratio + black + lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -15.5984  -2.7386  -0.5046   1.7273  26.2373
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  36.341145   5.067492   7.171 2.73e-12 ***
## crim        -0.108413   0.032779  -3.307 0.001010 **
## zn           0.045845   0.013523   3.390 0.000754 ***
## chas         2.718716   0.854240   3.183 0.001551 **
## nox        -17.376023   3.535243  -4.915 1.21e-06 ***
## rm           3.801579   0.406316   9.356 < 2e-16 ***
## dis         -1.492711   0.185731  -8.037 6.84e-15 ***
## rad           0.299608   0.063402   4.726 3.00e-06 ***
## tax          -0.011778   0.003372  -3.493 0.000521 ***
## ptratio     -0.946525   0.129066  -7.334 9.24e-13 ***
## black         0.009291   0.002674   3.475 0.000557 ***
## lstat       -0.522553   0.047424 -11.019 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.736 on 494 degrees of freedom
## Multiple R-squared:  0.7406, Adjusted R-squared:  0.7348
## F-statistic: 128.2 on 11 and 494 DF, p-value: < 2.2e-16
```